



Stefan Seman

ML Engineer | Data Scientist | PhD Researcher

Phone: +38162262131

Website: manse404.github.io/webpage

GitHub: github.com/ManSe404

ResearchGate: [ResearchGate](#)

Email: stefan.seman404@gmail.com

LinkedIn: linkedin.com/in/sseman

Scholar: [Google Scholar](#)

PROFESSIONAL SUMMARY

ML Engineer, Data Scientist, and PhD Researcher with a strong background in higher education and published research. Specialize in machine learning, AI, generative AI, and software development, with unique expertise bridging data science and biomedicine. Analytical thinker with a pragmatic problem-solving approach and excellent communication skills.

Core Focus: LLMs, RAG Systems, Software Development, Document Intelligence

TECHNICAL SKILLS

Languages

Python, SQL, Bash, LabVIEW

AWS Cloud

Lambda, S3, ECS, EC2, Step Functions, Bedrock

Databases

DynamoDB, PostgreSQL, Redis, Pinecone, Milvus, PgVector

Frameworks

FastAPI, Flask, PyTorch, PySpark, LangChain, LlamaIndex, Transformers

Azure Cloud

Data Factory, Synapse, DevOps, Functions, Azure Search, Azure OpenAI, Data Lake, Databricks

Tools

Docker, Git, Grafana, Jenkins

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer & Software Engineer

Oct 2024 - Present

404 Solutions

- Designed and operated multi-tenant AI systems for enterprise clients
- Delivered document intelligence solutions for automated extraction, classification, and summarization
- Built scalable cloud workflows on AWS (S3, Step Functions, Bedrock, APIs) supporting multiple organizations
- Deployed, monitored, and maintained production systems across distributed client environments
- Stack:** Python, AI Agents, FastAPI, RAG

Senior Data Scientist & Software Engineer in AI, R&D

Oct 2023 - Sept 2024

Publicis Sapient

- Led development from POC to fully production-grade systems
- Built enterprise RAG systems with retrieval optimization
- Delivered client-facing solutions and technical propositions
- Stack:** Python, FastAPI, LangChain, Transformers, PyTorch, Docker, Azure, KubeFlow, Redis, Milvus, AI Agents, Azure OpenAI

Data Scientist & Developer -> Senior Specialist

May 2022 - Oct 2023

Valcon Netherlands

- Led and mentored a team of 6 engineers
- Owned design of document processing microservices
- Built data extraction and document classification services from scratch
- **Stack:** Python, Azure DevOps, Azure Data Factory, Synapse, Data Lake, Docker

Data Scientist

Oct 2021 - Apr 2022

Optimal Systems GmbH

- Systematic model evaluation and hyperparameter tuning for ML classification
- Data-driven insights and testing/validation of ML systems
- **Stack:** Python, Docker, Machine Learning, Azure

Researcher

Mar 2020 - Mar 2022

BioIRC - Bioengineering R&D Center

- Built components for EU Horizon 2020 computational platform (SILICOFCM)
- ECG, CPET, and biomedical data processing; peer-reviewed publications
- **Stack:** Python, Statistical Analysis, SPSS, ECG, CPET, Biomedicine

PhD Researcher

Oct 2018 - Mar 2022

University of Belgrade

- Developed research methodology and statistical frameworks
- Applied machine learning to biomedical datasets; published papers internationally
- **Stack:** Python, PyTorch, SPSS, Machine Learning, Biostatistics

Sports Scientist

Oct 2014 - July 2018

Sport and Recreation Center Obliquus

- Built algorithms for signal processing and data collection
- Functional testing and performance analysis for athletes
- **Stack:** Python, SPSS, LabVIEW, Data Analysis

KEY PROJECTS

Full list: manse404.github.io/webpage/projects.html

AI-Powered Document Summarizer

RAG-based summarization pipeline with LLMs for unstructured corpora
Python, AWS Bedrock, S3, Step Functions, LangChain, LlamaIndex, RAG

AI-Code Copilot

Enterprise AI copilot for internal infrastructure and knowledge bases
Python, Azure OpenAI, LangChain, Milvus, Redis, Fine-Tuning, KubeFlow

Document Chatbot

Enterprise chatbot with optimized ingestion pipelines and retrieval
Python, FastAPI, Azure Search, Azure OpenAI, RAG, Docker

Document Data Extractor Service

Deep learning algorithms for extracting data from unstructured text
Python, Azure Data Factory, Synapse, PyTorch, Flask, Docker

Large-Scale ETL Pipeline

Serverless data ingestion and transformation on AWS
Python, AWS Lambda, S3, Step Functions, DynamoDB, ECS

SILICOFCM - EU Horizon 2020

Computational platform for in silico clinical trials of cardiomyopathies
silicofcm.eu
Python, PyTorch, Biostatistics, ECG, CPET

Document Classification Service

ML-based document categorization with containerized microservices
Python, Docker, Azure, Flask, REST API, ML Classification

Data Extraction Service

ETL platform with Azure Data Factory and Power BI integration
Python, FastAPI, Azure Data Factory, Synapse, Data Lake, Power BI

Voice-Enabled Support Chatbot

Conversational AI with speech-to-text and text-to-speech for telephony
Python, Azure LUIS, Azure Speech Services, Azure Bot Service, Docker

SELECTED PUBLICATIONS (TOTAL: 10+ PUBLICATIONS | 100+ CITATIONS)

A Machine Learning Approach to Gene Expression in Hypertrophic Cardiomyopathy

Pavic, J., Zivanovic, M.N., Seman, S., et al.
Pharmaceuticals (MDPI), 2024

Ventilatory Efficiency Parameters Outperform Peak VO₂ in Monitoring Therapy Effects in HCM

Seman, S., Tesic, M., Popovic, D., et al. (SILICOFCM Investigators)
European Heart Journal, Vol. 45, Supplement 1, 2024

Defining the Importance of Stress Reduction in Managing Cardiovascular Disease - The Role of Exercise

Popovic, D., Bjelobrk, M., Tesic, M., Seman, S., et al. (HL-PIVOT)
Progress in Cardiovascular Diseases, Vol. 70, pp. 84-93, 2022

Physical activity and exercise as an essential medical strategy for the COVID-19 pandemic and beyond

Seman, S., Srzentic Drazilov, S., et al. (HL-PIVOT Network)
Experimental Biology and Medicine, Vol. 246, Issue 21, 2021

EDUCATION

PHD CANDIDATE

Sport Science

University of Belgrade
2018 - Present

Focus: Machine learning applications in exercise physiology and biomedical signal processing

MASTER'S DEGREE

Sports Medicine

University of Novi Sad

BACHELOR'S DEGREE

Sport Science

University of Belgrade

CERTIFICATIONS

- ✓ Microsoft Certified: Azure AI Associate
- ✓ Databricks Academy: Lakehouse Fundamentals
- ✓ Johns Hopkins University: COVID-19 Contact Tracing (Coursera, 2021)
- ✓ Microsoft Certified: Azure Data Fundamentals
- ✓ Databricks Academy: Generative AI Fundamentals

MEMBERSHIPS

- ✓ ACSM - American College of Sports Medicine
- ✓ SEBM - Society for Experimental Biology and Medicine
- ✓ Serbian Society for Prevention of Chronic Noncommunicable Diseases
- ✓ Serbian Mountain Rescue Service

Languages: Serbian (Native) English (Fluent)